

LLVM Optimization Passes: A Study of Analysis and Transformation

Amaan Ali Sayed

1 Introduction

LLVM provides a powerful set of analysis and transformation passes designed to improve the intermediate representation (IR) used in compilers. This report explores three of these optimization passes in depth:

- **Mem2Reg** (Promote Memory to Register)
- **GVN** (Global Value Numbering)
- **Loop Unroll**

For each pass, we provide a description, illustrate its effect using a small Javalette program, and compare the original and optimized LLVM IR code. All code examples are compiled from a simple Javalette compiler and transformed using `opt` with the respective pass.

2 Mem2Reg (Promote Memory to Register)

Description

The Mem2Reg pass transforms memory-based variables (defined via `alloca`) into SSA-form register variables. It uses dominance frontiers to place `phi` functions where control flow converges, eliminating redundant memory accesses.

Javalette Program

```
int main() {  
    int x;  
    x = 42;  
    return x;  
}
```

LLVM IR Before Optimization (Generated by compiler)

```
define i32 @main() {  
entry:  
    %t0 = alloca i32  
    store i32 0, i32* %t0  
    store i32 42, i32* %t0
```

```
%t1 = load i32, i32* %t0
ret i32 %t1
}
```

LLVM IR After opt -mem2reg

```
define i32 @main() {
entry:
    ret i32 42
}
```

Explanation

The pass eliminates the stack allocation and replaces the load-store pair with a direct use of the constant value. This simplifies the IR and enables subsequent optimizations like constant folding.

3 GVN (Global Value Numbering)

Description

GVN identifies redundant expressions across the program and reuses their values. It assigns a unique number to each expression and eliminates those with equivalent computations.

Javalette Program

```
int main() {
    int a = 10;
    int b = 10;
    int c = a + b;
    int d = a + b;
    return c + d;
}
```

LLVM IR Before Optimization

```
define i32 @main() {
entry:
    %t0 = alloca i32
    store i32 10, i32* %t0
    %t1 = alloca i32
    store i32 10, i32* %t1
    %t3 = load i32, i32* %t0
    %t4 = load i32, i32* %t1
    %t5 = add i32 %t3, %t4
    %t2 = alloca i32
    store i32 %t5, i32* %t2
    %t7 = load i32, i32* %t0
    %t8 = load i32, i32* %t1
```

```

    %t9 = add i32 %t7, %t8
    %t6 = alloca i32
    store i32 %t9, i32* %t6
    %t10 = load i32, i32* %t2
    %t11 = load i32, i32* %t6
    %t12 = add i32 %t10, %t11
    ret i32 %t12
}

```

LLVM IR After opt -gvn

```

define i32 @main() {
entry:
    %t0 = alloca i32, align 4
    store i32 10, i32* %t0, align 4
    %t1 = alloca i32, align 4
    store i32 10, i32* %t1, align 4
    %t2 = alloca i32, align 4
    store i32 20, i32* %t2, align 4
    %t6 = alloca i32, align 4
    store i32 20, i32* %t6, align 4
    ret i32 40
}

```

Explanation

The pass detects that `a + b` is computed twice and eliminates the second computation. Furthermore, since both `a` and `b` are constants, constant propagation and folding are applied, resulting in a direct return of 40.

4 InstCombine (Instruction Combining)

Description

The InstCombine pass performs peephole optimizations on LLVM IR, combining sequences of instructions into simpler ones. It is not a canonicalization pass but rather focuses on local simplification, folding constants, and eliminating unnecessary computations. It often runs multiple times during optimization pipelines.

Javalette Program

```

int main() {
    int a = 3;
    int b = 4;
    int x = (a * 1) + (b * 0);
    return x;
}

```

LLVM IR Before Optimization

```
define i32 @main() {
entry:
    %t0 = alloca i32
    store i32 3, i32* %t0
    %t1 = alloca i32
    store i32 4, i32* %t1
    %t3 = load i32, i32* %t0
    %t4 = mul i32 %t3, 1
    %t5 = load i32, i32* %t1
    %t6 = mul i32 %t5, 0
    %t7 = add i32 %t4, %t6
    %t2 = alloca i32
    store i32 %t7, i32* %t2
    %t8 = load i32, i32* %t2
    ret i32 %t8
}
```

LLVM IR After opt -instcombine

```
define i32 @main() {
entry:
    ret i32 3
}
```

Explanation

The instruction combining pass applies the following simplifications:

- $\%a * 1 \rightarrow \%a$
- $\%b * 0 \rightarrow 0$
- $\%a + 0 \rightarrow \%a$

This reduces the entire computation to a constant value (3), allowing the function to return it directly. These small local simplifications make the IR more efficient and prepare it for further optimizations like constant propagation and dead code elimination.